

Deep Learning

An MIT Press book

Ian Goodfellow, Yoshua Bengio and Aaron Courville

[Exercises](#) [Lecture Slides](#)

The Deep Learning textbook is a resource intended to help students and practitioners enter the field of machine learning in general and deep learning in particular. The online version of the book is now complete and will remain available online for free. The print version will be available for sale soon.

Citing the book

To cite this book, please use this bibtex entry:

```
@unpublished{Goodfellow-et-al-2016-Book,  
  title={Deep Learning},  
  author={Ian Goodfellow, Yoshua Bengio, and Aaron Courville},  
  note={Book in preparation for MIT Press},  
  url={http://www.deeplearningbook.org},  
  year={2016}  
}
```

FAQ

- Can I get a PDF of this book?

No, our contract with MIT Press forbids distribution of too easily copied electronic formats of the book.

- Why are you using HTML format for the drafts?

This format is a sort of weak DRM required by our contract with MIT Press. It's intended to discourage unauthorized copying/editing of the book. Unfortunately, the conversion from PDF to HTML is not perfect, and some things like subscript expressions do not render correctly. If you have a suggestion for a better way of making the book available to a wide audience while preventing unauthorized copies, please let us know.

- What is the best way to print the HTML format?

Printing seems to work best printing directly from the browser, using Chrome. Other browsers do not work as well.

- When will the book come out?

It's difficult to predict. MIT Press is currently preparing the book for printing. Please contact us if you are interested in using the textbook for course materials in the short term; we will put you in contact with MIT Press.

If you notice any typos (besides the known issues listed below) or have suggestions for exercises to add to the website, do not hesitate to contact any of the authors directly by e-mail: Ian (<lastname.firstname>@gmail.com), Yoshua (<firstname>.<lastname>@umontreal.ca), Aaron (<firstname>.<lastname>@gmail.com). The book itself is now complete and we are not currently making revisions beyond correcting any minor errors that remain.

Known issues: The website currently does not display some symbols correctly, due to the conversion from PDF to HTML. Pending approval from MIT Press, we will change the format used for the website. Currently, the known issues are that in the Edge browser, the "does not equal" sign sometimes appears as the "equals" sign. This may be resolved by switching to Chrome. Also, some parentheses do not appear in any browser. For example, Eq. 5.29, 5.33, and 5.34 are missing parentheses in the HTML. There is currently no workaround for this issue other than to wait for us to receive approval to use another format.

- [Table of Contents](#)
- [Acknowledgements](#)
- [Notation](#)
- [1 Introduction](#)
- [Part I: Applied Math and Machine Learning Basics](#)
 - [2 Linear Algebra](#)
 - [3 Probability and Information Theory](#)
 - [4 Numerical Computation](#)
 - [5 Machine Learning Basics](#)
- [Part II: Modern Practical Deep Networks](#)
 - [6 Deep Feedforward Networks](#)
 - [7 Regularization](#)
 - [8 Optimization for Training Deep Models](#)
 - [9 Convolutional Networks](#)
 - [10 Sequence Modeling: Recurrent and Recursive Nets](#)
 - [11 Practical Methodology](#)
 - [12 Applications](#)

- [Part III: Deep Learning Research](#)
 - [13 Linear Factor Models](#)
 - [14 Autoencoders](#)
 - [15 Representation Learning](#)
 - [16 Structured Probabilistic Models for Deep Learning](#)
 - [17 Monte Carlo Methods](#)
 - [18 Confronting the Partition Function](#)
 - [19 Approximate Inference](#)
 - [20 Deep Generative Models](#)
- [Bibliography](#)
- [Index](#)